

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
!pip install shutup
```

```
In [ ]: import torch
import torchvision
import torch.nn as nn
import torchvision.transforms as transforms
import torchvision.datasets as dsets
from torchvision.datasets import MNIST
import matplotlib.pyplot as plt
import numpy as np
from torch.optim import Adam, SGD
import random
import math
import random
from random import choice
torch.manual_seed(17)
import os
from statistics import mode
from scipy import stats
import pickle
import cv2
import sklearn
from sklearn.model_selection import train_test_split
import time
import pandas as pd
import matplotlib
import shutup; shutup.please()
from torchsummary import summary
```

```
In [ ]: # def show_data(data_sample):
#     plt.imshow(data_sample[0].reshape(32,32), cmap='gray')
#     diff = np.max(np.absolute(np.subtract(data_sample[0],data_sample[3])))
#     plt.title('y_true = '+ str(data_sample[1]) + ' y_pred = '+str(data_sa

def show_data_2(data_sample, pred_label):
    plt.figure()
    plt.imshow(data_sample.reshape(28,28), cmap='gray')
    plt.title('y_pred = '+str(pred_label))
    plt.show()
```

```
In [ ]: x_train = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/ECE657A/A4/x
x_test = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/ECE657A/A4/x
y_train = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/ECE657A/A4/y
y_test = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/ECE657A/A4/y
```

```
In [ ]: x_train_data = x_train.values
x_test_data = x_test.values
y_train_data = y_train.values
y_test_data = y_test.values
```

```
In [ ]: x_train_data, x_val_data, y_train_data, y_val_data = train_test_split(x_t
```

```
In [ ]: def build_dataloader(data=x_train_data,label = y_train_data, batch_size =
        ip = []
        lab = []
        tmp_ip = []
        tmp_lab = []
        count = 0
        for t,l in zip(data,label):
            tmp_ip.append(cv2.resize(t,(32,32)).reshape(1,32,32))
            tmp_lab.append(int(l))
            count+=1
            if count%batch_size == 0:
                ip.append(torch.tensor(tmp_ip))
                lab.append(torch.tensor(np.array(tmp_lab)))
                tmp_ip = []
                tmp_lab = []
        return (ip,lab)

def preprocess_ip(x):
    return torch.tensor(cv2.resize(x,(32,32)).reshape(1,1,32,32))
```

```
In [ ]: loader= build_dataloader()
        val_loader = build_dataloader(data = x_val_data, label = y_val_data )
        test_loader = build_dataloader(data = x_test_data, label = y_test_data)
```

/usr/local/lib/python3.7/dist-packages/ipykernel\_launcher.py:12: UserWarning: Creating a tensor from a list of numpy.ndarrays is extremely slow. Please consider converting the list to a single numpy.ndarray with numpy.array() before converting to a tensor. (Triggered internally at ../torch/csrc/utils/tensor\_new.cpp:201.)

```
if sys.path[0] == '':
```

```
In [ ]: class CNN(nn.Module):
        def __init__(self):
            super(CNN, self).__init__()
            self.cnn1 = nn.Conv2d(in_channels = 1, out_channels = 32, kernel_size
            self.relu1 = nn.ReLU()
            self.maxpool1 = nn.MaxPool2d(kernel_size=2)
            self.cnn2 = nn.Conv2d(in_channels = 32, out_channels = 32, kernel_siz
            self.fc1 = nn.Linear(2048,5)
            self.softmax = nn.Softmax()

        def forward(self,x):
            out = self.cnn1(x)
            out = self.relu1(out)
            out = self.maxpool1(out)
            out = self.cnn2(out)
            out = self.relu1(out)
            out = self.maxpool1(out)
            out = out.view(out.size(0), -1)
            out = self.fc1(out)
            out = self.softmax(out)
            return out

        def train(self, train_loader,opt, loss_fn, BATCH_SIZE, eval = True):
            loss_tot = 0
            acc = 0
```

```

cnt = 0
for x,y in zip(train_loader[0], train_loader[1]):
    cnt+=1
    x = x.float()
    x = x.cuda()
    y = y.cuda()
    opt.zero_grad()
    z = self.forward(x)
    y = torch.nn.functional.one_hot(y, num_classes=5)
    loss = loss_fn(z.double(),y.double())
    loss_tot+= loss.item()
    loss.backward()
    opt.step()
    if cnt%1000==0:
        print("Loss is ",loss.item())
loss_tr = loss_tot/(len(train_loader[1]))
if eval:
    correct = 0
    for x,y in zip(train_loader[0],train_loader[1]):
        x = x.float()
        x = x.cuda()
        y = y.cuda()
        z = self.forward(x)
        _,yhat = torch.max(z.data,1)
        correct += (yhat== y).sum().item()
    acc = correct/(len(train_loader[1])*BATCH_SIZE)
return loss_tr,acc

def eval(self,loader,BATCH_SIZE,loss_fn=nn.CrossEntropyLoss()):
    correct = 0
    acc = 0
    loss_tot = 0
    for x,y in zip(loader[0],loader[1]):
        x = x.float()
        x = x.cuda()
        y = y.cuda()
        z = self.forward(x)
        loss = loss_fn(z.double(),y.long())
        loss_tot += loss.item()
        _, yhat = torch.max(z.data,1)
        correct += (yhat == y).sum().item()
    acc = correct/(len(loader[1]))
    loss_val = loss_tot/(len(loader[1])*BATCH_SIZE)
    return loss_val, acc

```

## [CM1] Classification with CNN

### Default Network

#### Model 1 - summary

```
In [ ]: model = CNN()  
model = model.cuda()  
summary(model, (1,32,32))
```

```
-----  
              Layer (type)              Output Shape              Param #  
=====
```

Conv2d-1	[-1, 32, 32, 32]	320
ReLU-2	[-1, 32, 32, 32]	0
MaxPool2d-3	[-1, 32, 16, 16]	0
Conv2d-4	[-1, 32, 16, 16]	9,248
ReLU-5	[-1, 32, 16, 16]	0
MaxPool2d-6	[-1, 32, 8, 8]	0
Linear-7	[-1, 5]	10,245
Softmax-8	[-1, 5]	0

```
=====
```

Total params: 19,813  
Trainable params: 19,813  
Non-trainable params: 0

```
-----
```

Input size (MB): 0.00  
Forward/backward pass size (MB): 0.70  
Params size (MB): 0.08  
Estimated Total Size (MB): 0.78

```
-----
```

**Network Architecture:**

We add convolutional layers and flatten the final result to feed into the densely connected layers. Finally we add the densely connected layers. We configure the network as per the given network specification.

**Activation Functions:****1. Relu**

effectively means "If  $X > 0$  return  $X$ , else return 0" -- so what it does is it only passes values 0 or greater to the next layer in the network. ReLU stands for Rectified Linear Unit. The main advantage of using the ReLU function over other activation functions is that it does not activate all the neurons at the same time. This means that the neurons will only be deactivated if the output of the linear transformation is less than 0. For the negative input values, the result is zero, that means the neuron does not get activated. Since only a certain number of neurons are activated, the ReLU function is far more computationally efficient when compared to the sigmoid and tanh function. Due to this reason, during the backpropagation process, the weights and biases for some neurons are not updated.

**2. Softmax**

Since our problem is a classification problem, we employ softmax as the activation function of last layer.

**Optimizer:** We have used "adam" optimizer while compiling the model because Adam combines the best properties of the AdaGrad and RMSProp algorithms to provide an optimization algorithm that can handle sparse gradients on noisy problems. Adam is relatively easy to configure where the default configuration parameters do well on most problems.


**Loss Function:** Since our problem is a classification problem, we use CrossEntropyLoss.

```
In [ ]: opt = Adam(model.parameters(), lr = 0.001)
        crit = nn.CrossEntropyLoss()
        model = model.cuda()

        tr_loss = []
        tr_acc = []
        val_acc = []
        val_loss = []
        training_time_tot = 0
        for epoch in range(50):
            start_time = time.time()
            l, a = model.train(train_loader=loader, opt = opt, loss_fn = crit, BATCH
            end_time = time.time()
            training_time_tot += end_time - start_time
            print("Epoch ", epoch+1, "--> ", a)
            tr_loss.append(l)
            tr_acc.append(a)
            l, a = model.eval(loader=val_loader, BATCH_SIZE=100)
            val_acc.append(a)
            val_loss.append(l)
```


Epoch	1 -->	0.793
Epoch	2 -->	0.8001458333333333
Epoch	3 -->	0.8012291666666667
Epoch	4 -->	0.8219791666666667
Epoch	5 -->	0.828625
Epoch	6 -->	0.8310416666666667
Epoch	7 -->	0.8275
Epoch	8 -->	0.8319791666666667
Epoch	9 -->	0.8355
Epoch	10 -->	0.8405416666666666
Epoch	11 -->	0.8321875
Epoch	12 -->	0.8409583333333334
Epoch	13 -->	0.8414583333333333
Epoch	14 -->	0.8387916666666667
Epoch	15 -->	0.8468958333333333
Epoch	16 -->	0.8455208333333334
Epoch	17 -->	0.8419375
Epoch	18 -->	0.8466041666666667
Epoch	19 -->	0.8512916666666667
Epoch	20 -->	0.84575
Epoch	21 -->	0.8535416666666666
Epoch	22 -->	0.852
Epoch	23 -->	0.8468958333333333
Epoch	24 -->	0.8460208333333333
Epoch	25 -->	0.8494375
Epoch	26 -->	0.8529791666666666
Epoch	27 -->	0.8526458333333333
Epoch	28 -->	0.8529791666666666
Epoch	29 -->	0.8561875
Epoch	30 -->	0.8540208333333333
Epoch	31 -->	0.8549583333333334
Epoch	32 -->	0.8591041666666667
Epoch	33 -->	0.8511666666666666
Epoch	34 -->	0.8551666666666666
Epoch	35 -->	0.8571875
Epoch	36 -->	0.8565416666666666
Epoch	37 -->	0.850875
Epoch	38 -->	0.8475416666666666
Epoch	39 -->	0.8569791666666666
Epoch	40 -->	0.8555833333333334
Epoch	41 -->	0.8559583333333334
Epoch	42 -->	0.8534166666666667
Epoch	43 -->	0.8575625
Epoch	44 -->	0.8529375
Epoch	45 -->	0.8560416666666667
Epoch	46 -->	0.8602083333333334
Epoch	47 -->	0.8563333333333333
Epoch	48 -->	0.86275
Epoch	49 -->	0.861125
Epoch	50 -->	0.8459791666666666

```
In [ ]: plt.plot(tr_acc, label = 'Training Acc')
plt.plot(np.array(val_acc)/100.0, label = 'Validation Acc')
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.title("Accuracy by Epoch plot - Model 1")
plt.legend()
plt.show()
```



```
In [ ]: plt.plot(np.array(tr_loss)/100, label = 'Training Loss')
plt.plot(np.array(val_loss), label = 'Validation Loss')
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Loss by Epoch plot - Model 1")
plt.legend()
```

Out [ ]: <matplotlib.legend.Legend at 0x7ff48b218d90>





```
In [ ]: test_loss, test_acc = model.eval(loader=test_loader, BATCH_SIZE=100)
        print("Test loss for Model 1 is ", test_loss)
        print("Test accuracy for Model 1 is ", test_acc)
```

```
Test loss for Model 1 is  0.010652118769183854
```

```
Test accuracy for Model 1 is  83.85
```

### **Model 1 Performance:**

Based on training and validation accuracy, we can see that there is no overfitting that is affecting the model as validation accuracy closely follows training accuracy. We conclude training of the model after 50 epochs. Test accuracy for this model after training for 50 epochs is 83.85%. The factors upon which the accuracy depends includes a lot of factors outside of network architecture. Some of those factors are learning rate, number of epochs, network initialization etc.